

Hand On with R

Task 1: Basic building

R can be used as an interactive calculator

```
> 5 + 7
```

```
[1] 12
```

Similarly

```
x <- 5 + 7
```

Now let's create a vector

```
z <- c(1.1, 9, 3.4)
```

Any time you can get help using

```
> ?c
```

combine vectors to make a new vector

```
c(z, 555, z)
```

```
[1] 1.10  9.00  3.14 555.00  1.10  9.00  3.14
```

Numeric vectors can be used in arithmetic expressions

```
> z * 2 + 100
```

```
[1] 102.20 118.00 106.28
```

Other common arithmetic operators are ``+``, ``-``, ``/``, and ``^`` (where `x^2` means 'x squared'). To take the square root, use the `sqrt()` function and to take the absolute value, use the `abs()` function.

For example :

```
> my_sqrt <- sqrt(z - 1)
```

Will return a vector of length 3

```
> my_sqrt
```

```
[1] 0.3162278 2.8284271 1.4628739
```

Recycling (broadcasting in python) in R

```
> c(1, 2, 3, 4) + c(0, 10)
```

```
[1] 1 12 3 14
```

However,

```
> c(1, 2, 3, 4) + c(0, 10, 100)
```

```
[1] 1 12 103 4
```

This will give you a warning message

Warning message:

In $c(1, 2, 3, 4) + c(0, 10, 100)$:

longer object length is not a multiple of shorter object length

Tip:

The up arrow will cycle through previous commands.

Task 2: Workspace and Files

In this task, you'll learn how to examine your local workspace in R and begin to explore the relationship between your workspace and the file system of your machine.

```
> getwd()
```

```
[1] "C:/Users/Nitin PC/Documents"
```

Will give you directory your R session is using as its current working directory using `getwd()`.

List all the objects in your local workspace using `ls()`.

```
> ls()
```

```
[1] "x", "y", "z", "my_sqrt"
```

List all the files in your working directory using `list.files()` or `dir()`

Try it yourself

```
> list.files()
```

What is the output?

Let's Play

```
> old.dir<- getwd()
```

Use `dir.create()` to create a directory in the current working directory called "testdir".

```
> dir.create("testdir")
```

Set your working directory to "testdir" with the `setwd()` command.

```
> setwd("testdir")
```

You can create a file in your working directory called "mytest.R" using the `file.create()` function.

```
> file.create("mytest.R")
```

```
[1] TRUE
```

Let's validate:

```
> list.files()
```

```
[1] "mytest.R"
```

Or

```
> file.exists("mytest.R")
```

```
[1] TRUE
```

Now when an object or a file is created it has certain attributes which can be accessed by

```
> file.info("mytest.R")
```

What is the output ?

You can use the \$ operator --- e.g., `file.info("mytest.R")$mode` --- to grab specific items.

Renaming a file

```
> file.rename("mytest.R", "mytest2.R")
```

```
[1] TRUE
```

Coping a file

```
> file.copy("mytest2.R", "mytest3.R")
```

```
[1] TRUE
```

Go back to your original working directory using `setwd()`.

```
> setwd(old.dir)
```

Task 3: Sequences of Numbers

The simplest way to create a sequence of numbers in R is by using the `:` operator.

```
> 1:20
```

```
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
```

Try

```
> pi:10
```

```
[1] 3.141593 4.141593 5.141593 6.141593 7.141593 8.141593 9.141593
```

Try

```
> 15:1
```

```
[1] 15 14 13 12 11 10 9 8 7 6 5 4 3 2 1
```

It counted backwards in increments of 1!

Pull up the documentation for `:` now.

```
> ?':'
```

You could also use

```
> seq(0, 10, by=0.5)
```

```
[1] 0.0 0.5 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5 6.0 6.5 7.0 7.5
```

```
[17] 8.0 8.5 9.0 9.5 10.0
```

```
> seq(5, 10, length=30)
```

```
[1] 5.000000 5.172414 5.344828 5.517241 5.689655 5.862069 6.034483 6.206897
```

```
[9] 6.379310 6.551724 6.724138 6.896552 7.068966 7.241379 7.413793 7.586207
```

```
[17] 7.758621 7.931034 8.103448 8.275862 8.448276 8.620690 8.793103 8.965517
```

```
[25] 9.137931 9.310345 9.482759 9.655172 9.827586 10.000000
```

If we're interested in creating a vector that contains 40 zeros, we can use `rep(0, times = 40)`. Try it out.

```
> rep(0, times=40)
```

```
[1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

Give another try instead type

```
> rep(c(0, 1, 2), times = 10)
```

```
[1] 0 1 2 0 1 2 0 1 2 0 1 2 0 1 2 0 1 2 0 1 2 0 1 2 0 1 2 0 1 2 0 1 2
```

Task 4: Vectors

The simplest and most common data structure in R is the vector(Array in C , C++ , JAVA)

Two types:

- 1) Atomic (exactly one data type)
- 2) List (multiple data types)

Types of atomic vectors include logical, character, integer, and complex.

```
> num_vect <- c(0.5, 55, -10, 6)
```

Try

```
> num_vect < 1
```

Will return a vector with 4 logical values

```
[1] TRUE FALSE TRUE FALSE
```

Logical operators

The ``<`` and ``>=`` symbols in these examples are called 'logical operators'. Other logical operators include ``>``, ``<=``, ``==`` for exact equality, and ``!=`` for inequality.

Other logical operators are

If we have two logical expressions, A and B, we can ask whether at least one is TRUE with A | B (logical 'or' a.k.a. 'union') or whether they are both TRUE with A & B (logical 'and' a.k.a. 'intersection'). Lastly, !A is the negation of A and is TRUE when A is FALSE and vice versa.

```
((111 >= 111) | !(TRUE)) & ((4 + 1) == 5)
```

Will return TRUE

Create a character vector that contains the following words: "Programming", "with", "R".

```
> my_char <- c("Programming", "with", "R")
```

```
> my_char
```

```
[1] "Programming", "with", "R"
```

Type `paste(my_char, collapse = " ")` now. Make sure there's a space between the double quotes in the ``collapse`` argument.

```
> paste(my_char, collapse = " ")
```

```
[1] "Programming with R"
```

The ``collapse`` argument to the `paste()` function tells R that when we join together the elements of the `my_char` character vector, we'd like to separate them with single spaces.

Vector recycling! Try `paste(LETTERS, 1:4, sep = "-")`, where `LETTERS` is a predefined variable in R containing a character vector of all 26 letters in the English alphabet.

```
> paste(LETTERS, 1:4, sep = "-")
```

[1] "A-1" "B-2" "C-3" "D-4" "E-1" "F-2" "G-3" "H-4" "I-1" "J-2" "K-3" "L-4" "M-1" "N-2"

[15] "O-3" "P-4" "Q-1" "R-2" "S-3" "T-4" "U-1" "V-2" "W-3" "X-4" "Y-1" "Z-2"

Task 5: Missing Values

Missing values play an important role in statistics and data analysis. Often, missing values must not be ignored, but rather they should be carefully studied to see if there's an underlying pattern or cause for their missingness.

In R, NA is used to represent any value that is 'not available' or 'missing' (in the statistical sense).

```
> x <- c(44, NA, 5, NA)
```

```
> x * 3
```

```
[1] 132 NA 15 NA
```

Next, let's create a vector containing 1000 NAs with `z <- rep(NA, 1000)`.

```
> z <- rep(NA, 1000)
```

Finally, let's select 100 elements at random from these 2000 values (combining `y` and `z`) such that we don't know how many NAs we'll wind up with or what positions they'll occupy in our final vector -- `my_data <- sample(c(y, z), 100)`.

```
> my_data <- sample(c(y, z), 100)
```

```
> my_na <- is.na(my_data)
```

```
> my_na
```

```
[1] TRUE TRUE TRUE FALSE TRUE FALSE TRUE TRUE TRUE FALSE FALSE FALSE TRUE
```

```
[14] FALSE FALSE TRUE FALSE TRUE FALSE TRUE TRUE FALSE FALSE FALSE TRUE FALSE
```

```
[27] TRUE FALSE TRUE TRUE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
[40] FALSE TRUE TRUE FALSE FALSE TRUE FALSE FALSE TRUE FALSE TRUE FALSE FALSE
```

```
[53] TRUE TRUE TRUE TRUE FALSE FALSE TRUE TRUE FALSE TRUE TRUE TRUE TRUE
```

```
[66] TRUE FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE FALSE FALSE
```

```
[79] FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE
```

```
[92] FALSE TRUE TRUE FALSE FALSE TRUE TRUE TRUE FALSE
```

Let's give that a try here. Call the `sum()` function on `my_na` to count the total number of TRUEs in `my_na`, and thus the total number of NAs in `my_data`. Don't assign the result to a new variable.

```
> sum(my_na)
```

```
[1] 52
```

Let's look at a second type of missing value -- NaN, which stands for 'not a number'. To generate NaN, try dividing (using a forward slash) 0 by 0 now.

```
> 0 / 0
```

```
[1] NaN
```


Task 6 : Subsetting Vectors

Extract elements from a vector based on some conditions that we specify.

Try `x[1:10]` to view the first ten elements of `x`.

```
> x[1:10]
```

```
[1]      NA 0.5299247274      NA      NA -0.5399988732 -0.1191854427
[7]      NA -0.0005441663 0.9714911664      NA
```

Recall that `!` gives us the negation of a logical expression, so `!is.na(x)` can be read as 'is not NA'. Therefore, if we want to create a vector called `y` that contains all of the non-NA values from `x`, we can use `y <- x[!is.na(x)]`. Give it a try.

```
> y <- x[!is.na(x)]
```

```
> y
```

```
[1] 0.5299247274 -0.5399988732 -0.1191854427 -0.0005441663 0.9714911664 -1.0325497198
[7] 0.3798177474 -0.2863554538 -0.3914337301 -0.5790076778 0.4994686757 -0.3207320342
[13] -0.2557526850 1.7753843380 -0.7728033356 0.5825444776 -0.6870112749 1.4382822638
[19] -1.7819660619 1.2992963698
```

Now type `y[y > 0]` to see that we get all of the positive elements of `y`, which are also the positive elements of our original vector `x`.

```
> y[y > 0]
```

```
[1] 0.5299247 0.9714912 0.3798177 0.4994687 1.7753843 0.5825445 1.4382823 1.2992964
```

```
> x[x > 0]
```

```
[1]      NA 0.5299247      NA      NA      NA 0.9714912      NA      NA
[9]      NA      NA 0.3798177      NA      NA      NA      NA 0.4994687
[17]      NA      NA      NA      NA      NA 1.7753843 0.5825445      NA
[25] 1.4382823 1.2992964      NA      NA
```

Since NA is not a value, but rather a placeholder for an unknown quantity, the expression `NA > 0` evaluates to NA. Hence we get a bunch of NAs mixed in with our positive numbers when we do this.

We could do this:

```
> x[!is.na(x) & x > 0]
```

```
[1] 0.5299247 0.9714912 0.3798177 0.4994687 1.7753843 0.5825445 1.4382823 1.2992964
```

Create a vector of indexes with `c(3, 5, 7)`, then put that inside of the square brackets.

```
> x[c(3, 5, 7)]
```

```
[1]      NA -0.5399989      NA
```

Try `x[c(-2, -10)]` gives us all elements of `x` EXCEPT for the 2nd and 10 elements.

```
> x[c(-2, -10)]
```

```
[1]      NA      NA      NA -0.5399988732 -0.1191854427      NA
[7] -0.0005441663  0.9714911664      NA      NA      NA -1.0325497198
[13] 0.3798177474      NA -0.2863554538      NA      NA -0.3914337301
[19]      NA -0.5790076778  0.4994686757      NA      NA      NA
[25] -0.3207320342 -0.2557526850      NA      NA  1.7753843380 -0.7728033356
[31] 0.5825444776      NA -0.6870112749  1.4382822638 -1.7819660619  1.2992963698
[37]      NA      NA
```

Task 7: Matrices and Data Frames

Create a vector

```
> my_vector <- 1:20
```

The `dim()` function tells us the 'dimensions' of an object.

```
> dim(my_vector)
```

```
NULL
```

```
> length(my_vector)
```

```
[1] 20
```

We can assign dimension of a vector

```
> dim(my_vector) <- c(4, 5)
```

Another way to see dimensions is by calling the `attributes()` function on `my_vector`. Try it now.

```
> attributes(my_vector)
```

```
$dim
```

```
[1] 4 5
```

With more than one dimension it is no more a vector rather it is a Matrix

```
> my_vector
```

```
 [,1] [,2] [,3] [,4] [,5]
```

```
[1,]  1   5   9  13  17
```

```
[2,]  2   6  10  14  18
```

```
[3,]  3   7  11  15  19
```

```
[4,]  4   8  12  16  20
```

You can verify the same

```
> class(my_vector)
```

```
[1] "matrix"
```

Another way of defining a matrix is

```
> my_matrix2 <- as.matrix(1:20)
```

```
> my_matrix2
```

```
 [,1]
```

```
[1,]  1
```

```
[2,]  2
```

```
[3,]  3
```

```
[4,] 4
[5,] 5
[6,] 6
[7,] 7
[8,] 8
[9,] 9
[10,] 10
[11,] 11
[12,] 12
[13,] 13
[14,] 14
[15,] 15
[16,] 16
[17,] 17
[18,] 18
[19,] 19
[20,] 20
```

```
> matrix(1:20, nrow = 4, ncol = 5, byrow = FALSE)
     [,1] [,2] [,3] [,4] [,5]
[1,]  1   5   9  13  17
[2,]  2   6  10  14  18
[3,]  3   7  11  15  19
[4,]  4   8  12  16  20
```

Let's start by creating a character vector containing the names of our patients -- Bill, Gina, Kelly, and Sean. Remember that double quotes tell R that something is a character string. Store the result in a variable called patients.

```
> patients <- c("Bill", "Gina", "Kelly", "Sean")
```

Now we'll use the cbind() function to 'combine columns'. Don't worry about storing the result in a new variable. Just call cbind() with two arguments -- the patients vector and my_matrix.

```
> cbind(patients, my_matrix)

patients
[1,] "Bill" "1" "5" "9" "13" "17"
[2,] "Gina" "2" "6" "10" "14" "18"
[3,] "Kelly" "3" "7" "11" "15" "19"
[4,] "Sean" "4" "8" "12" "16" "20"
```

All the data is in quotes or as strings. So in order to have more than one data type we have data frames

```
> my_data <- data.frame(patients, my_matrix)
```

```
> my_data
  patients X1 X2 X3 X4 X5
1   Bill  1  5  9 13 17
2   Gina  2  6 10 14 18
3  Kelly  3  7 11 15 19
4   Sean  4  8 12 16 20
```

```
1 Bill 1 5 9 13 17
2 Gina 2 6 10 14 18
3 Kelly 3 7 11 15 19
4 Sean 4 8 12 16 20
```

Verify

```
Class(my_data)
```

```
[1] "data.frame"
```

Task 8: Logic

This lesson is meant to be an introduction to logical operations in R

Two logical value

you can construct logical expressions which will evaluate to either TRUE or FALSE

The first logical operator we are going to discuss is the equality operator, represented by two equals signs `==`. Use the equality operator below to find out if TRUE is equal to TRUE.

```
TRUE == TRUE
```

```
[1] TRUE
```

Try

```
(FALSE == TRUE) == FALSE
```

We can compare numbers also

```
> 6 == 7
```

```
[1] FALSE
```

Similarly

```
> 6 > 7
```

Other operators

`<=`, `>=`, `!=`, etc

Which of the following evaluates to FALSE?

1: `10 == 10`

2: `0 > -34`

3: `7 >= 11`

4: `60 < 80`

Not equal to operator in R `!=`

```
> 6 != 7
```

Which of the following evaluates to FALSE?

1: `!(0 >= -1)`

2: `!FALSE`

3: $9 < 10$

4: $7 \neq 8$

What do you think the following expression will evaluate to?: $(\text{TRUE} \neq \text{FALSE}) == !(6 == 7)$

1: FALSE

2: TRUE

To examine relationship between multiple logical expression we can use AND and OR operators

AND - `&` and `&&` (both are right in R)

You can use the `&` operator to evaluate AND across a vector. The `&&` version of AND only evaluates the first member of a vector. Let's test both for practice. Type the expression `TRUE & c(TRUE, FALSE, FALSE)`.

Try

```
> TRUE & c(TRUE, FALSE, FALSE)
```

```
[1] TRUE FALSE FALSE
```

What happens in this case is that the left operand `TRUE` is recycled across every element in the vector of the right operand. This is the equivalent statement as `c(TRUE, TRUE, TRUE) & c(TRUE, FALSE, FALSE)`.

Then we use `&&`

```
> TRUE && c(TRUE, FALSE, FALSE)
```

```
[1] TRUE
```

| In this case, the left operand is only evaluated with the first member of the right operand (the vector). The rest of the elements in the vector aren't evaluated at all in this expression.

The OR operator (`|`)

The `|` version of OR evaluates OR across an entire vector, while the `||` version of OR only evaluates the first member of a vector.

```
> TRUE | c(TRUE, FALSE, FALSE)
```

```
[1] TRUE TRUE TRUE
```

Now let's try out the non-vectorized version of the OR operator. Type the expression `TRUE || c(TRUE, FALSE, FALSE)`.

```
> TRUE || c(TRUE, FALSE, FALSE)
```

```
[1] TRUE
```

Remember the order of operations in arithmetic operation likewise AND operations are evaluated before the OR operator

Now try

```
5 > 8 || 6 != 8 && 4 > 3.9
```

Let try another one

Which one of the following expressions evaluates to TRUE?

1: FALSE || TRUE && FALSE

2: TRUE && 62 < 62 && 44 >= 44

3: 99.99 > 100 || 45 < 7.3 || 4 != 4.0

4: TRUE && FALSE || 9 >= 4 && 3 < 6

Let take advantage of some of the R in built function

isTRUE() take one argument

>isTRUE(6 > 4)

Which of the following evaluates to TRUE?

1: !isTRUE(4 < 3)

2: isTRUE(!TRUE)

3: !isTRUE(8 != 5)

4: isTRUE(3)

5: isTRUE(NA)

The function identical() will return TRUE if the two R objects passed to it as arguments are identical. Try out

identical('ABC', 'ABC')

Which of the following evaluates to TRUE?

1: identical('hello', 'Hello')

2: identical(5 > 4, 3 < 3.1)

3: !identical(7, 7)

4: identical(4, 3.1)

Other functions

xor()

Lets create a vector

>mod <- sample(10)

Display

>mod

mod is a random sampling of integers from 1 to 10 without replacement.

Try

```
mod>5
```

what is the output?

A vector of logical values

Now we can use this logical vector to ask other questions

which() functions

so

Example: `which(c(TRUE,FALSE,TRUE))`

Use the `which()` function to find the indices of mod that are greater than 7

Which of the following commands would produce the indices of the elements in mod that are less than or equal to 4?

1: `mod <= 4`

2: `which(mod <= 4)`

3: `which(mod < 4)`

4: `mod < 4`

`any()` and `all()` take logical vectors as their argument. The `any()` function will return TRUE if one or more of the elements in the logical vector is TRUE. The `all()` function will return TRUE if every element in the logical vector is TRUE.

Try

```
>any(mod<0)
```

Which of the following evaluates to TRUE?

1: `all(c(TRUE, FALSE, TRUE))`

2: `any(mod == 4.5)`

3: `any(mod == 20)`

4: `all(mod == 10)`